

# HPC-19-2

February 21, 2024

## 1 Assignment 11 (Canon's Matrix Multiplication)

1. Describe Canon's Matrix Multiplication algorithm.

A. Cannon's algorithm is a distributed matrix multiplication algorithm that works particularly well with two-dimensional meshes. Lynn Elliot Cannon described it for the first time in 1969.

Algorithm Overview:

- \* When multiplying two  $n \times n$  matrices A and B, we need  $n \times n$  processing nodes arranged in a 2D.
- \* Each processing node (PE) performs the following steps:
  1. Initialize:  $k = (i + j) \bmod n$  (to ensure different processors access distinct data).
  2. Compute:  $c[i][j] += a[i][k] \times b[k][j]$  concurrently.
  3. Communicate:
    - \* Send a to PE  $(i, (j + N - 1) \bmod n)$ .
    - \* Send b to PE  $((i + N - 1) \bmod n, j)$ .
    - \* Receive a' from PE  $(i, (j + 1) \bmod n)$ .
    - \* Receive b' from PE  $((i + 1) \bmod n, j)$ .
  4. Update:  $a = a'$  and  $b = b'$ .
- \* Repeat the above steps for  $n$  iterations, ensuring that each processor accesses different matrix elements.
- \* The storage requirements remain constant and independent of the number of processors.

2. Implement Canon's Matrix Multiplication using collective communication.

3. Analyze the efficiency of the code.

```
[1]: from mpi4py import MPI
import numpy as np
```

```
[2]: comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()
```

```
[3]: import time
import matplotlib.pyplot as plt
```

```
[4]: N = 2000
```

```

[5]: if N % size != 0:
        raise ValueError("Matrix size N must be divisible by the number of_
        ↪processes (size)")

    block_size = N // size

    print(f"Rank {rank}: Starting execution")

    if rank == 0:
        print(f"Rank {rank}: Generating matrices A and B")
        A = np.random.randint(0, 10, (N, N))
        B = np.random.randint(0, 10, (N, N))
        print(f"Rank {rank}: Matrices A and B generated")
    else:
        A = None
        B = None

    print(f"Rank {rank}: Broadcasting matrices A and B")

```

```

Rank 0: Starting execution
Rank 0: Generating matrices A and B
Rank 0: Matrices A and B generated
Rank 0: Broadcasting matrices A and B

```

```

[6]: start_time = time.time()
    A = comm.bcast(A, root=0)
    B = comm.bcast(B, root=0)
    end_time = time.time()
    print(f"Rank {rank}: Matrices A and B broadcasted")

```

```

Rank 0: Matrices A and B broadcasted

```

```

[7]: A_rows = np.zeros((block_size, N), dtype=int)
    comm.Scatter(A, A_rows, root=0)

    start_time_multiplication = time.time()
    C_rows = np.dot(A_rows, B)
    end_time_multiplication = time.time()

    C = None
    if rank == 0:
        C = np.zeros((N, N), dtype=int)

    comm.Gather(C_rows, C, root=0)

    if rank == 0:
        print("Resultant Matrix C:")
        print(C)

```

```

print("Broadcasting time:", end_time - start_time, "seconds")
print("Matrix multiplication time:", end_time_multiplication -
↪start_time_multiplication, "seconds")

```

Resultant Matrix C:

```

[[40747 41349 41354 ... 41260 41169 39328]
 [39955 41279 41428 ... 40482 42471 40155]
 [40196 41163 40919 ... 40635 39846 39888]
 ...
 [40674 41117 40256 ... 40439 39831 39460]
 [39447 40800 40591 ... 40467 40723 40124]
 [40330 40734 40834 ... 41727 41303 40117]]

```

Broadcasting time: 0.039728403091430664 seconds

Matrix multiplication time: 20.929131031036377 seconds

[8]: `!mpirun -n 4 python mpi_scatter_gather.py`

```

Rank 3: Starting execution
Rank 3: Broadcasting matrices A and B
Rank 3: Matrices A and B broadcasted
Rank 1: Starting execution
Rank 1: Broadcasting matrices A and B
Rank 1: Matrices A and B broadcasted
Rank 2: Starting execution
Rank 2: Broadcasting matrices A and B
Rank 2: Matrices A and B broadcasted
Rank 0: Starting execution
Rank 0: Generating matrices A and B
Rank 0: Matrices A and B generated
Rank 0: Broadcasting matrices A and B
Rank 0: Matrices A and B broadcasted

```

Resultant Matrix C:

```

[[41048 39762 40853 ... 41543 41611 39830]
 [40264 38648 41718 ... 40252 40400 39149]
 [41054 40522 41801 ... 41724 41400 41007]
 ...
 [41452 39388 41118 ... 41178 40805 40147]
 [39720 38856 40107 ... 40178 40473 39127]
 [40218 38532 40695 ... 40379 40558 39088]]

```

Broadcasting time: 0.03373289108276367 seconds

Matrix multiplication time: 4.518505096435547 seconds

[9]: `!mpirun -n 8 python mpi_scatter_gather.py`

```

Rank 5: Starting execution
Rank 5: Broadcasting matrices A and B
Rank 5: Matrices A and B broadcasted
Rank 7: Starting execution
Rank 7: Broadcasting matrices A and B

```

```

Rank 7: Matrices A and B broadcasted
Rank 6: Starting execution
Rank 6: Broadcasting matrices A and B
Rank 6: Matrices A and B broadcasted
Rank 1: Starting execution
Rank 1: Broadcasting matrices A and B
Rank 1: Matrices A and B broadcasted
Rank 3: Starting execution
Rank 3: Broadcasting matrices A and B
Rank 3: Matrices A and B broadcasted
Rank 2: Starting execution
Rank 2: Broadcasting matrices A and B
Rank 2: Matrices A and B broadcasted
Rank 4: Starting execution
Rank 4: Broadcasting matrices A and B
Rank 4: Matrices A and B broadcasted
Rank 0: Starting execution
Rank 0: Generating matrices A and B
Rank 0: Matrices A and B generated
Rank 0: Broadcasting matrices A and B
Rank 0: Matrices A and B broadcasted
Resultant Matrix C:
[[39642 40083 40950 ... 41662 41105 40729]
 [40628 40791 40643 ... 41148 41905 39657]
 [40344 41142 40593 ... 42146 42763 40476]
 ...
 [39234 39860 40460 ... 40807 40930 40160]
 [38715 39396 39736 ... 41051 39867 39587]
 [40204 40537 40572 ... 41915 41841 40698]]
Broadcasting time: 0.066436767578125 seconds
Matrix multiplication time: 2.638921022415161 seconds

```

```
[10]: !mpirun -n 16 python mpi_scatter_gather.py
```

```

Rank 5: Starting execution
Rank 5: Broadcasting matrices A and B
Rank 5: Matrices A and B broadcasted
Rank 15: Starting execution
Rank 15: Broadcasting matrices A and B
Rank 15: Matrices A and B broadcasted
Rank 9: Starting execution
Rank 9: Broadcasting matrices A and B
Rank 9: Matrices A and B broadcasted
Rank 11: Starting execution
Rank 11: Broadcasting matrices A and B
Rank 11: Matrices A and B broadcasted
Rank 10: Starting execution
Rank 10: Broadcasting matrices A and B

```

```

Rank 10: Matrices A and B broadcasted
Rank 1: Starting execution
Rank 1: Broadcasting matrices A and B
Rank 1: Matrices A and B broadcasted
Rank 3: Starting execution
Rank 3: Broadcasting matrices A and B
Rank 3: Matrices A and B broadcasted
Rank 2: Starting execution
Rank 2: Broadcasting matrices A and B
Rank 2: Matrices A and B broadcasted
Rank 7: Starting execution
Rank 7: Broadcasting matrices A and B
Rank 7: Matrices A and B broadcasted
Rank 6: Starting execution
Rank 6: Broadcasting matrices A and B
Rank 6: Matrices A and B broadcasted
Rank 4: Starting execution
Rank 4: Broadcasting matrices A and B
Rank 4: Matrices A and B broadcasted
Rank 13: Starting execution
Rank 13: Broadcasting matrices A and B
Rank 13: Matrices A and B broadcasted
Rank 14: Starting execution
Rank 14: Broadcasting matrices A and B
Rank 14: Matrices A and B broadcasted
Rank 12: Starting execution
Rank 12: Broadcasting matrices A and B
Rank 12: Matrices A and B broadcasted
Rank 8: Starting execution
Rank 8: Broadcasting matrices A and B
Rank 8: Matrices A and B broadcasted
Rank 0: Starting execution
Rank 0: Generating matrices A and B
Rank 0: Matrices A and B generated
Rank 0: Broadcasting matrices A and B
Rank 0: Matrices A and B broadcasted
Resultant Matrix C:
[[39895 40041 38242 ... 39939 39545 40169]
 [39345 39674 38612 ... 39800 40560 38977]
 [40013 39937 38095 ... 40518 41096 40771]
 ...
 [40008 39913 39012 ... 39942 41684 39822]
 [41780 41378 40286 ... 41451 42157 41894]
 [40301 40239 39599 ... 39864 40844 40837]]
Broadcasting time: 0.22203373908996582 seconds
Matrix multiplication time: 3.5083491802215576 seconds

```

```

[11]: N = 2000
      if N % size != 0:
          raise ValueError("Matrix size N must be divisible by the number of_
          ↳processes (size)")

      block_size = N // size

      print(f"Rank {rank}: Starting execution")

      if rank == 0:
          print(f"Rank {rank}: Generating matrices A and B")
          A = np.random.randint(0, 10, (N, N))
          B = np.random.randint(0, 10, (N, N))
          print(f"Rank {rank}: Matrices A and B generated")
      else:
          A = None
          B = None

      print(f"Rank {rank}: Broadcasting matrices A and B")

```

```

Rank 0: Starting execution
Rank 0: Generating matrices A and B
Rank 0: Matrices A and B generated
Rank 0: Broadcasting matrices A and B

```

```

[12]: start_time = time.time()
      A = comm.bcast(A, root=0)
      B = comm.bcast(B, root=0)
      end_time = time.time()
      print(f"Rank {rank}: Matrices A and B broadcasted")

      A_rows = np.zeros((block_size, N), dtype=int)
      comm.Scatter(A, A_rows, root=0)

      start_time_multiplication = time.time()
      C_rows = np.dot(A_rows, B)
      end_time_multiplication = time.time()

      start_time_gather = time.time()
      C_all = np.zeros((N, N), dtype=int)
      comm.Allgather(C_rows, C_all)
      end_time_gather = time.time()

      if rank == 0:
          print("Resultant Matrix C:")
          print(C_all)
          print("Broadcasting time:", end_time - start_time, "seconds")

```

```

print("Gathering time:", end_time_gather - start_time_gather, "seconds")
print("Matrix multiplication time:", end_time_multiplication -
↪start_time_multiplication, "seconds")

```

```

Rank 0: Matrices A and B broadcasted
Resultant Matrix C:
[[40262 41419 41149 ... 40627 41185 39118]
 [40512 40635 39829 ... 40602 41313 39061]
 [41364 40627 40990 ... 40016 41141 40741]
 ...
 [40798 40859 41008 ... 39674 40921 40082]
 [40201 40859 40919 ... 40100 40550 39726]
 [41429 40674 40897 ... 39753 41346 39610]]
Broadcasting time: 0.036681175231933594 seconds
Gathering time: 0.007696866989135742 seconds
Matrix multiplication time: 21.168152570724487 seconds

```

```
[13]: !mpirun -n 4 python MPI_Allgather.py
```

```

Rank 3: Starting execution
Rank 3: Broadcasting matrices A and B
Rank 3: Matrices A and B broadcasted
Rank 1: Starting execution
Rank 1: Broadcasting matrices A and B
Rank 1: Matrices A and B broadcasted
Rank 2: Starting execution
Rank 2: Broadcasting matrices A and B
Rank 2: Matrices A and B broadcasted
Rank 0: Starting execution
Rank 0: Generating matrices A and B
Rank 0: Matrices A and B generated
Rank 0: Broadcasting matrices A and B
Rank 0: Matrices A and B broadcasted
Resultant Matrix C:
[[39268 39927 39869 ... 39620 39577 39546]
 [40260 41715 40449 ... 41332 40204 40137]
 [39132 40005 39775 ... 40137 39549 40140]
 ...
 [40205 40717 40398 ... 40964 40263 40094]
 [40544 41437 40248 ... 42028 40126 40540]
 [40018 40157 40004 ... 41190 39445 40415]]
Broadcasting time: 0.031401872634887695 seconds
Gathering time: 0.011379480361938477 seconds
Matrix multiplication time: 4.781134605407715 seconds

```

```
[14]: !mpirun -n 8 python MPI_Allgather.py
```

```

Rank 3: Starting execution
Rank 3: Broadcasting matrices A and B

```

```

Rank 3: Matrices A and B broadcasted
Rank 2: Starting execution
Rank 2: Broadcasting matrices A and B
Rank 2: Matrices A and B broadcasted
Rank 4: Starting execution
Rank 4: Broadcasting matrices A and B
Rank 4: Matrices A and B broadcasted
Rank 7: Starting execution
Rank 7: Broadcasting matrices A and B
Rank 7: Matrices A and B broadcasted
Rank 1: Starting execution
Rank 1: Broadcasting matrices A and B
Rank 1: Matrices A and B broadcasted
Rank 0: Starting execution
Rank 0: Generating matrices A and B
Rank 0: Matrices A and B generated
Rank 0: Broadcasting matrices A and B
Rank 0: Matrices A and B broadcasted
Resultant Matrix C:
[[40351 39324 40644 ... 40178 40722 38738]
 [41633 41139 41735 ... 41205 41598 40359]
 [42168 40904 41892 ... 40341 42459 40579]
 ...
 [41025 40337 40735 ... 41337 41649 40040]
 [39452 38465 39429 ... 39067 39286 38344]
 [41148 40563 41671 ... 41142 41364 40395]]
Broadcasting time: 0.06987953186035156 seconds
Gathering time: 0.05123305320739746 seconds
Matrix multiplication time: 2.8613476753234863 seconds
Rank 5: Starting execution
Rank 5: Broadcasting matrices A and B
Rank 5: Matrices A and B broadcasted
Rank 6: Starting execution
Rank 6: Broadcasting matrices A and B
Rank 6: Matrices A and B broadcasted

```

```
[15]: !mpirun -n 16 python MPI_Allgather.py
```

```

Rank 9: Starting execution
Rank 9: Broadcasting matrices A and B
Rank 9: Matrices A and B broadcasted
Rank 10: Starting execution
Rank 10: Broadcasting matrices A and B
Rank 10: Matrices A and B broadcasted
Rank 7: Starting execution
Rank 7: Broadcasting matrices A and B
Rank 7: Matrices A and B broadcasted
Rank 8: Starting execution

```



Rank 8: Broadcasting matrices A and B  
 Rank 8: Matrices A and B broadcasted  
 Rank 5: Starting execution  
 Rank 5: Broadcasting matrices A and B  
 Rank 5: Matrices A and B broadcasted  
 Rank 6: Starting execution  
 Rank 6: Broadcasting matrices A and B  
 Rank 6: Matrices A and B broadcasted  
 Rank 11: Starting execution  
 Rank 11: Broadcasting matrices A and B  
 Rank 11: Matrices A and B broadcasted  
 Rank 13: Starting execution  
 Rank 13: Broadcasting matrices A and B  
 Rank 13: Matrices A and B broadcasted  
 Rank 12: Starting execution  
 Rank 12: Broadcasting matrices A and B  
 Rank 12: Matrices A and B broadcasted  
 Rank 0: Starting execution  
 Rank 0: Generating matrices A and B  
 Rank 0: Matrices A and B generated  
 Rank 0: Broadcasting matrices A and B  
 Rank 0: Matrices A and B broadcasted  
 Resultant Matrix C:  
 [[40098 40513 39370 ... 40675 40116 40372]  
 [40373 40444 39136 ... 40684 39919 40297]  
 [40017 40324 39607 ... 39830 40199 39921]  
 ...  
 [40336 40603 39674 ... 40866 39942 40451]  
 [40586 40233 39441 ... 40566 39837 40573]  
 [40127 41018 40006 ... 41142 40468 40251]]  
 Broadcasting time: 0.11317896842956543 seconds  
 Gathering time: 0.7809598445892334 seconds  
 Matrix multiplication time: 2.607423782348633 seconds  
 Rank 14: Starting execution  
 Rank 14: Broadcasting matrices A and B  
 Rank 14: Matrices A and B broadcasted  
 Rank 15: Starting execution  
 Rank 15: Broadcasting matrices A and B  
 Rank 15: Matrices A and B broadcasted  
 Rank 1: Starting execution  
 Rank 1: Broadcasting matrices A and B  
 Rank 1: Matrices A and B broadcasted  
 Rank 2: Starting execution  
 Rank 2: Broadcasting matrices A and B  
 Rank 2: Matrices A and B broadcasted  
 Rank 4: Starting execution  
 Rank 4: Broadcasting matrices A and B  
 Rank 4: Matrices A and B broadcasted

Rank 3: Starting execution  
Rank 3: Broadcasting matrices A and B  
Rank 3: Matrices A and B broadcasted

```
[22]: scatter_gather_processes = [4, 8, 16]
scatter_gather_broadcasting_time = [0.03373289108276367, 0.066436767578125, 0.
↳22203373908996582]
scatter_gather_multiplication_time = [4.518505096435547, 2.638921022415161, 3.
↳5083491802215576]

allgather_processes = [4, 8, 16]
allgather_broadcasting_time = [0.031401872634887695, 0.06987953186035156, 0.
↳11317896842956543]
allgather_multiplication_time = [4.781134605407715 , 2.8613476753234863, 2.
↳607423782348633]

plt.figure(figsize=(10, 5))

plt.subplot(1, 2, 1)
plt.plot(scatter_gather_processes, scatter_gather_broadcasting_time,
↳marker='o', label='MPI Scatter Gather')
plt.plot(allgather_processes, allgather_broadcasting_time, marker='o',
↳label='MPI All gather')
plt.xlabel('Number of Processes')
plt.ylabel('Broadcasting Time (seconds)')
plt.title('Broadcasting Time Comparison')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(scatter_gather_processes, scatter_gather_multiplication_time,
↳marker='o', label='MPI Scatter Gather')
plt.plot(allgather_processes, allgather_multiplication_time, marker='o',
↳label='MPI All gather')
plt.xlabel('Number of Processes')
plt.ylabel('Matrix Multiplication Time (seconds)')
plt.title('Matrix Multiplication Time Comparison')
plt.legend()

plt.tight_layout()
plt.show()
```

